



AWS Fargate

Team Members: Leonardo Egidati, Baptist Krugler, Matteo Salari, Xavier Lopez

Course: Cloud Computing & Big Data Analytics

Date: 12th May 2025

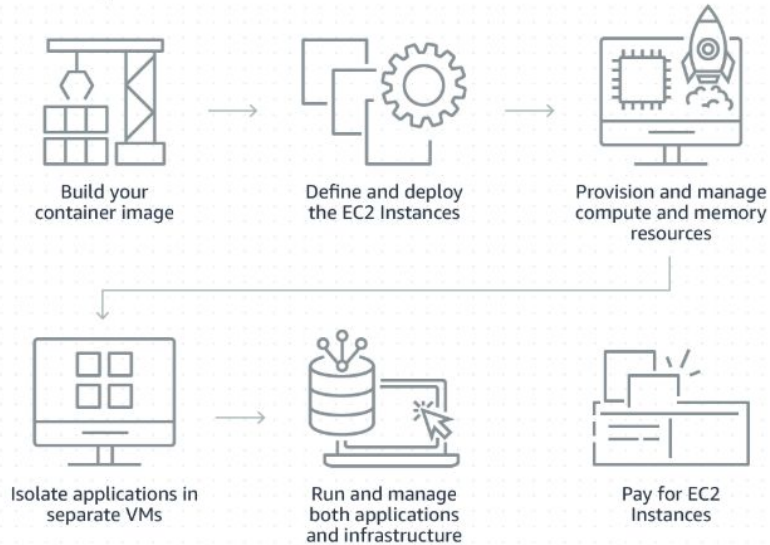
What is AWS Fargate?

A serverless compute engine for containers on AWS

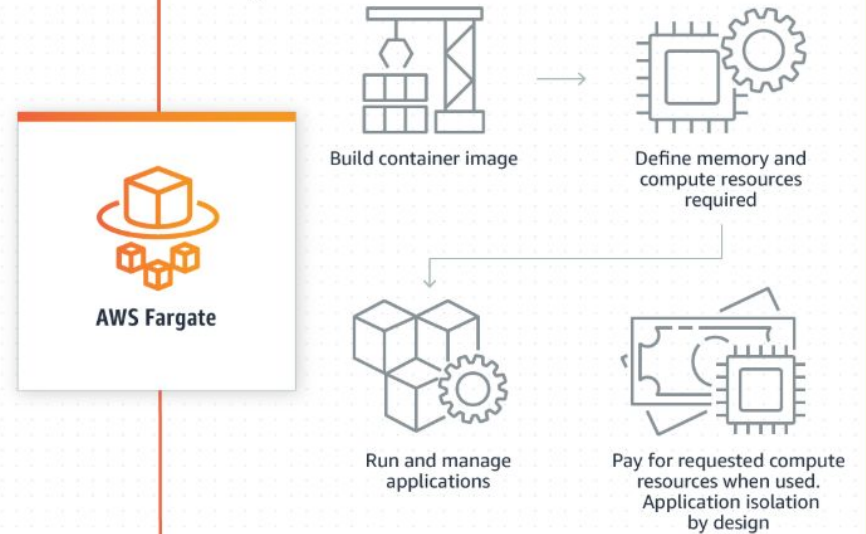
- **No servers to manage** – forget EC2 nodes or Kubernetes worker sets.
- **Works with ECS & EKS** – choose AWS's native orchestrator (ECS) or managed Kubernetes (EKS).
- **Declarative resources** – you specify vCPU, memory, IAM role, and networking; Fargate provisions the rest.
- **Per-task micro-VM isolation** – every task / pod runs in its own lightweight VM for strong security boundaries.
- **Automatic scaling & patching** – availability, OS updates, and cluster capacity are handled by AWS.
- **Pay for what you use** – per-second billing on CPU + RAM, no charge when containers aren't running.

You build the container, AWS runs it.

Without Fargate



With Fargate

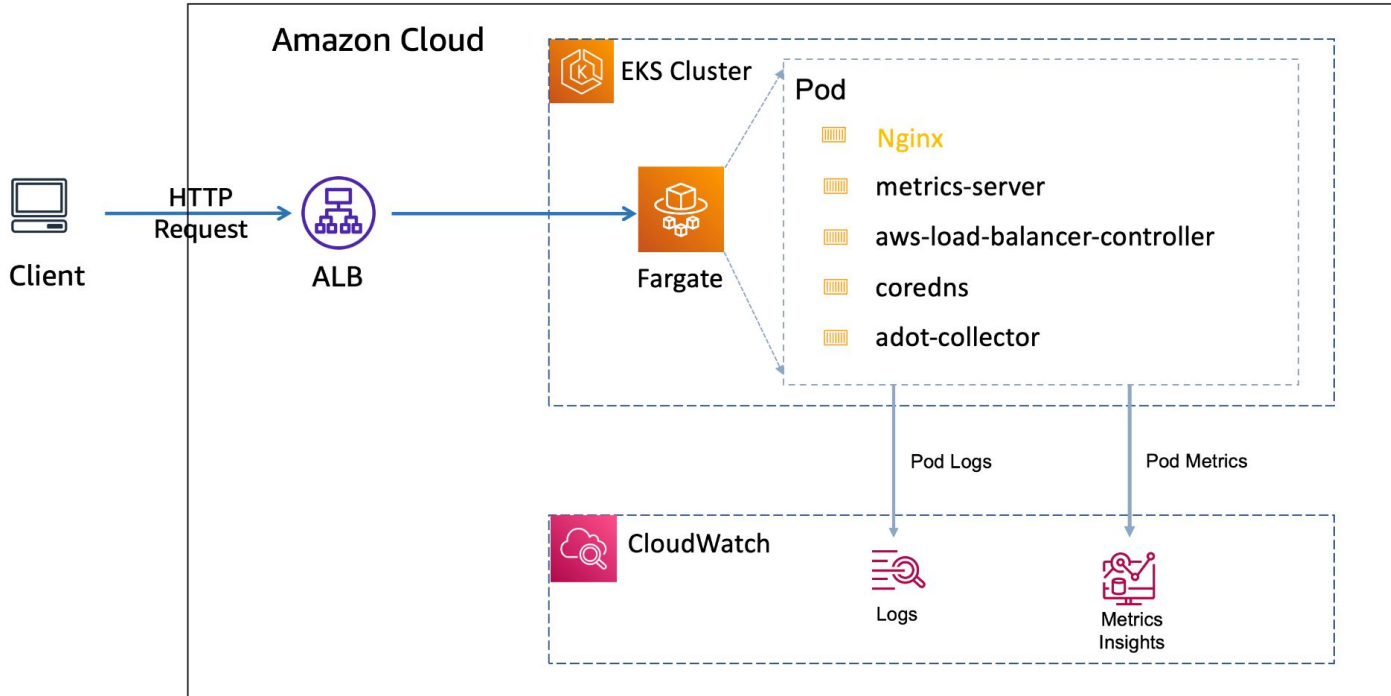


When to use it

Ideal scenarios

- **Spiky or unpredictable traffic**
Auto-scales in seconds without pre-provisioning capacity (e.g., flash-sale website, event-driven APIs).
- **Microservices & SaaS**
Dozens of small, independent services where managing EC2 fleets would be overhead.
- **Batch & data pipelines**
Short-lived ETL, log processing, or event-driven jobs—pay only while containers run.
- **Secure multi-tenant apps**
Each tenant's task runs in its own micro-VM, reducing “noisy-neighbor” risk.
- **CI/CD preview environments**
Spin up disposable review stacks on every pull request, tear them down after tests.
- **Regulated workloads needing host abstraction**
Isolation handled by AWS, simplifying compliance checks (PCI-DSS, HIPAA, etc.).

Basic architecture





Quick comparison

AWS Fargate

Serverless: no EC2 hosts to provision, patch, or scale.

Granular billing: pay per second for the vCPU + RAM your tasks use.

Automatic scaling: tasks spin up or down with load, even to zero.

Strong isolation: each task/pod runs inside its own Firecracker micro-VM.

Low ops overhead: you declare image, CPU, memory, IAM role, and run.

EC2 / Self-managed Kubernetes

Full host control: pick instance types, install agents, tweak kernel.

Reserved/Spot savings: cheaper for steady 24×7 workloads when hosts stay busy.

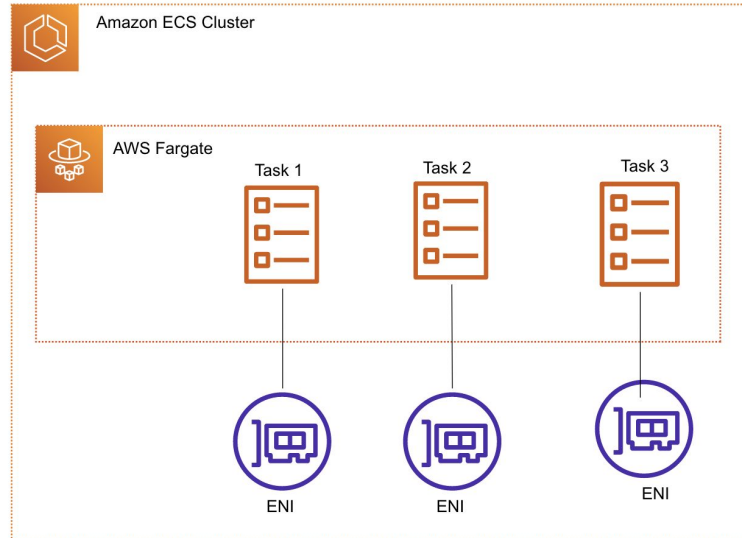
Manual capacity management: you size, scale, and patch the cluster.

Manual security management: you need to configure isolation and security manually.

Higher ops burden: upgrades, security fixes, autoscaling logic are on you.

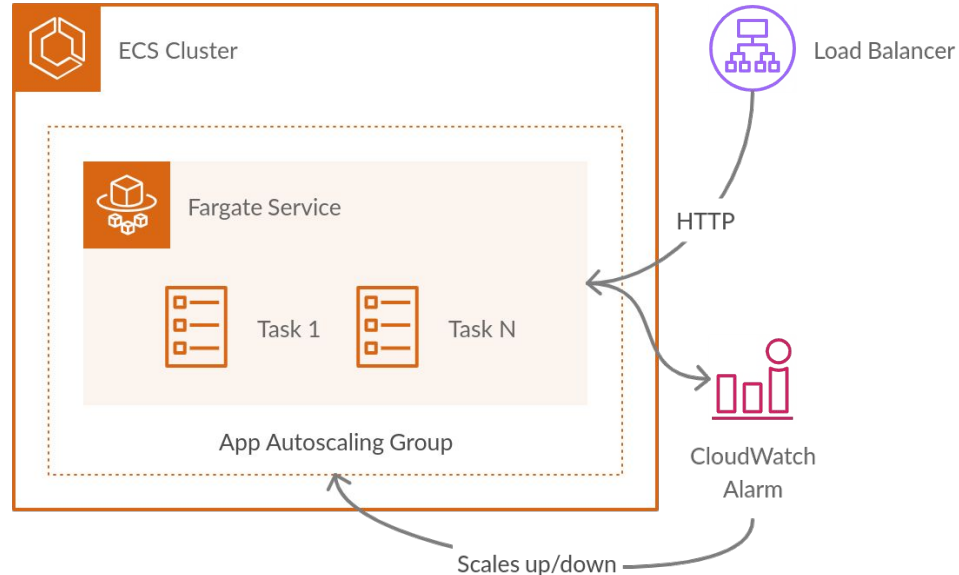
Part 1: Single task on Fargate

Goal: Understand the basic components of AWS Fargate in ECS. Define an ECS task, run a task in a cluster and understand the benefits of running a service instead of a task.



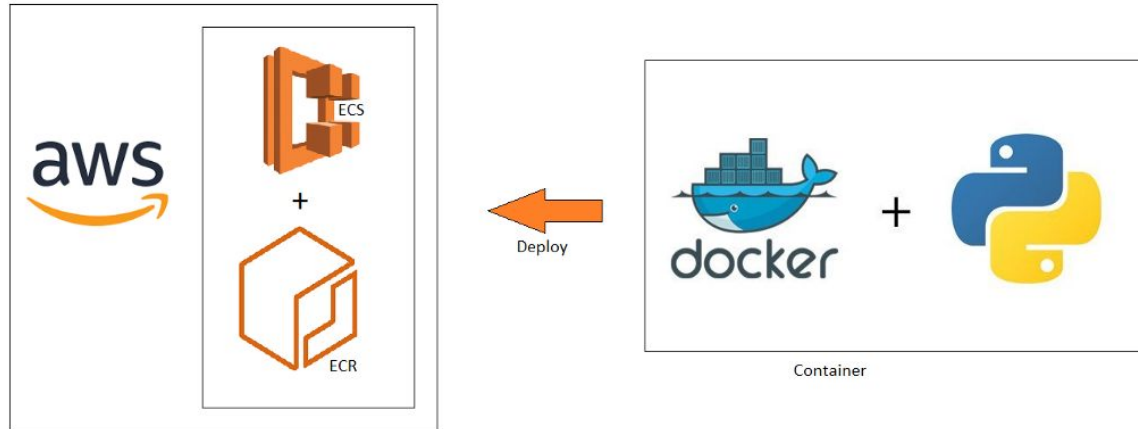
Part 2: Fargate+ALB+Auto-Scaling

Goal: Deploy a containerized web application using AWS Fargate and ECS, expose it via an Application Load Balancer (ALB), configure automatic scaling based on CPU usage, and safely clean up all resources.



Part 3: Custom image pipeline

Goal: Running a custom image in AWS ECS with Fargate. Upload a custom image into AWS ECR and defining a new task with or image. Understand how a cluster can have different services, one with each custom image.



Key takeaways & QA

Key points to remember

- **Serverless for containers** – Fargate removes host management; you focus on code.
- **Elastic by design** – tasks scale up, down, or to zero automatically.
- **Pay-only-when-running** – per-second billing on vCPU + RAM, no idle cost.
- **Strong isolation** – each task/pod gets its own micro-VM for security.
- **Best fit** – bursty microservices, short-lived jobs, rapid CI/CD previews.
- **Watch-outs** – higher cost for 24×7 loads, no GPU, limited host tweaks.
- **Decision rule** – choose Fargate when agility beats deep infrastructure control.