# AWS SageMaker

## CCBDA 2025

Aguirre Guaman, Pablo Fabricio  Bestak, Vojtech  Bång, Valdemar
Eberhardt, Jakob  Kosinski, Bartosz

May 12, 2025

# Agenda

- Everything-as-a-Service for Machine Learning
- Tutorial
- Conclusion

# Everything-as-a-Service for Machine Learning

- AWS SageMaker started in 2017 [1]
- Started as a platform for analytics, optimization, and machine learning
- Renamed to AWS SageMaker AI in 2024 [2]
- **Focus on the machine learning workflow**

## Data

- Fully integrated into AWS data management services, e.g. EMR [3]
- Data labeling and model validation AWS Ground Truth [4] and Mechanical Turk [5]
- Cloud services for human workforce

# Everything-as-a-Service for Machine Learning

## Development

- Jupyter notebooks for preprocessing and development
- SageMaker has a widely used Python SDK [6]
- And a IDE called SageMaker Studio [7]

# Everything-as-a-Service for Machine Learning

## Deployment

- Run pre-built or self-developed models on accelerated EC2 instances
- The more responsive, the more expensive [8]
- E.g. real-time inference with auto-scaling vs. serverless vs. batch transform overnight

# Tutorial

- Predict shared bike availability
- Load data into the cloud
- Define and train the model
- Provision it securely via an endpoint

# IAM Configuration

- ↑ Permissions for SageMaker resources are handled through IAM
- We need to grant the user access to SM-related resources
- ↓ We recommend doing it in a real (paid) AWS account to make sure all resources are available

```
"Statement": [
  {
    "Effect": "Allow",
    "Principal": {
      "Service": "sagemaker.amazonaws.com"
    },
    "Action": "sts:AssumeRole"
  }
]
```

**Figure:** Create a new IAM Role in **IAM** → **Access Management** → **Roles** called `AmazonSageMaker-TrainingExecutionRole`

# S3 Configuration

- ↑ SM can get data from all kinds of services and formats
- ↑ We settled for S3 because it is the simplest solution for testing
- ↓ Other persistent resources (e.g. RDS) can be expensive, even if we turn them off

Create new **S3 bucket**

- Set name to `ccbda-research-sagemaker`
- Leave all other settings as default

# Create Jupyter Notebook for development

- Jupyter notebooks are great for exploring and plotting data
- ↑ Run seamless on AWS
- ↑ Easy to collaborate with others
- ↓ Run on remote accelerated EC2 instances which can become expensive
- ↑ Users can download and execute notebooks step-by-step, e.g. our tutorial in `sagemaker_ml.ipynb`

# Handling the dataset

We are gonna use the **Seoul Bike Sharing Demand** dataset [9].

↑ Each cell in the notebook corresponds to:

- Loading the dataset
- Cleaning the data
- Splitting into train and test set
- Uploading it to our S3 bucket

# Training Job Configuration

```
xgboost_image_uri = image_uris.retrieve("xgboost",
↪  region=region, version="1.5-1")
estimator = Estimator(
    image_uri=xgboost_image_uri,
    role=role,
    instance_count=1,
    instance_type="ml.m5.large",
    volume_size=5,
    ...
)
estimator.set_hyperparameters(
    objective="reg:squarederror",
    num_round=100,
    max_depth=5,
    subsample=0.8,
    ...
)
```

# Training Job goes to the Queue

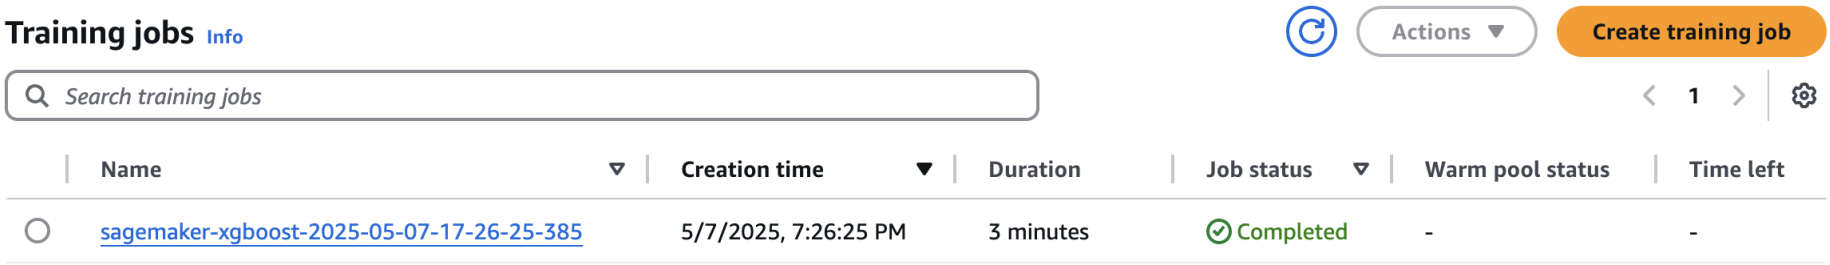


**Figure:** Created training job

# CloudWatch Integration

- ↑ Training and deployment logs are integrated in CloudWatch

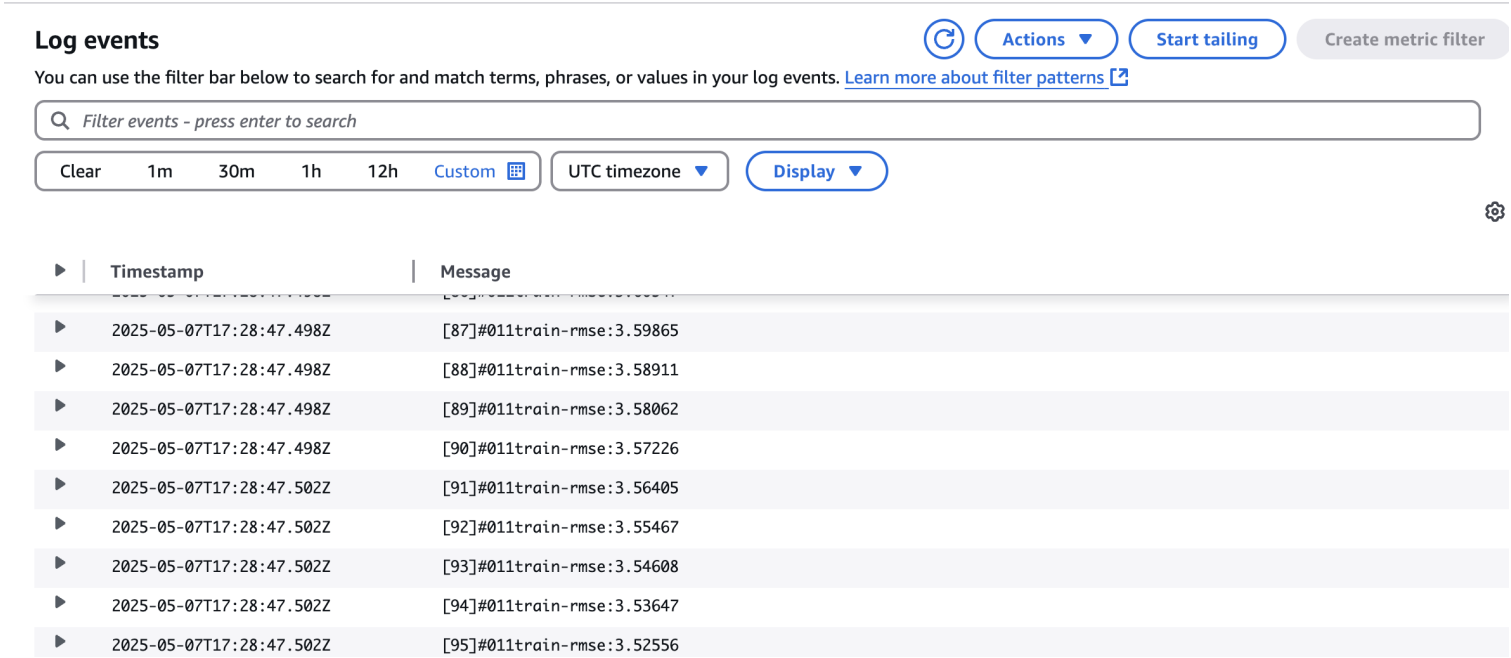


**Figure:** Training Job Logs

# Deploying Endpoint from the Notebook

- ↑ With a single line, we can deploy the trained model
- ↑ We have full control from here, e.g. about instance size, scaling policy

```
predictor = estimator.deploy(initial_instance_count=1,
↪  instance_type="ml.m5.large")
```

# The Endpoint is deployed & scaled automatically

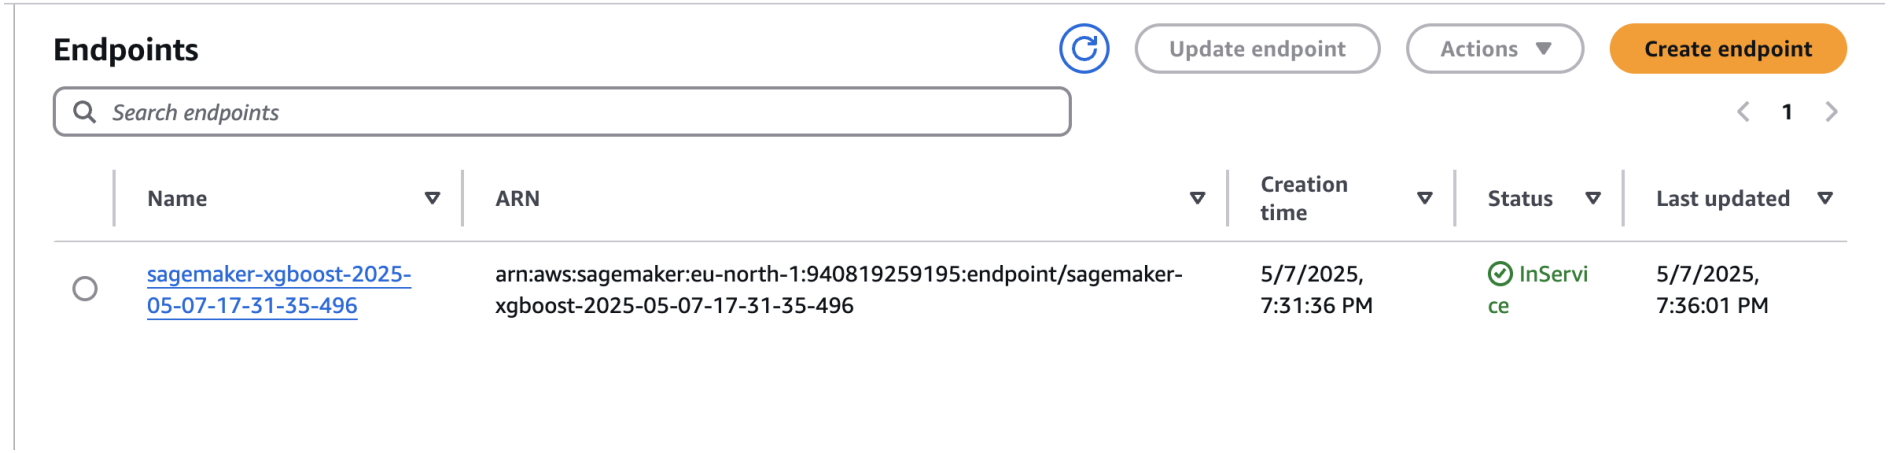


**Figure:** Published endpoint

# Conclusion & Opinion

- ↑ AWS SageMaker AI is feature complete for all aspects of ML
- ↑ We got started very fast thanks to the notebooks
- ↑ We have fine grain control over expected performance and QoS
- ↑↓ AWS provides adequate tooling, e.g. Python SDK
- ↑ Features industry-standard libraries, e.g. TensorFlow

# References & Questions I


[1] Ron Miller. *AWS releases SageMaker to make it easier to build and deploy machine learning models*. 8 May 2025. 2017. URL: https://techcrunch.com/2017/11/29/aws-releases-sagemaker-to-make-it-easier-to-build-and-deploy-machine-learning-models/.

[2] AWS. *Amazon SageMaker AI Rename*. https://docs.aws.amazon.com/sagemaker/latest/dg/whatis.html. Accessed: 2025-05-08. 2025.

[3] Amazon Web Services. *Amazon EMR*. https://aws.amazon.com/emr/. Accessed: 2025-05-08.

[4] Amazon Web Services. *SageMaker Ground Truth*. https://aws.amazon.com/sagemaker-ai/groundtruth/. Accessed: 2025-05-08.

# References & Questions II


[5] Amazon Web Services. *Using the Amazon Mechanical Turk Workforce*. https://docs.aws.amazon.com/sagemaker/latest/dg/sms-workforce-management-public.html. Accessed: 2025-05-08.

[6] Amazon Web Services. *Amazon SageMaker Notebook Instances and SDK*. https://docs.aws.amazon.com/sagemaker/latest/dg/nbi.html. Accessed: 2025-05-08.

[7] Amazon Web Services. *Amazon SageMaker Studio*. https://aws.amazon.com/sagemaker-ai/studio/?nc1=h_ls. Accessed: 2025-05-08.

[8] Amazon Web Services. *Model deployment options in Amazon SageMaker AI*. https://docs.aws.amazon.com/sagemaker/latest/dg/how-it-works-deployment.html. Accessed: 2025-05-08.

# References & Questions III

[9] UCI Machine Learning Repository. *Seoul Bike Sharing Demand*. https://archive.ics.uci.edu/dataset/560/seoul+bike+sharing+demand. Accessed: 2025-05-05. 2020. DOI: 10.24432/C5F62R.